



UNIVERSITÀ  
DEGLI STUDI  
FIRENZE

# Mathematical Morphology

## Parallel Computing Mid-Term

Angelo D'Amante  
Fabian Greavu




# Introduction

## Introduction: MM Theory

**Mathematical Morphology (MM) is a theory and technique for the analysis and processing of geometrical structures.**

**MM is also the foundation of morphological image processing, which consists of a set of operators that transform images according to any characterizations.**

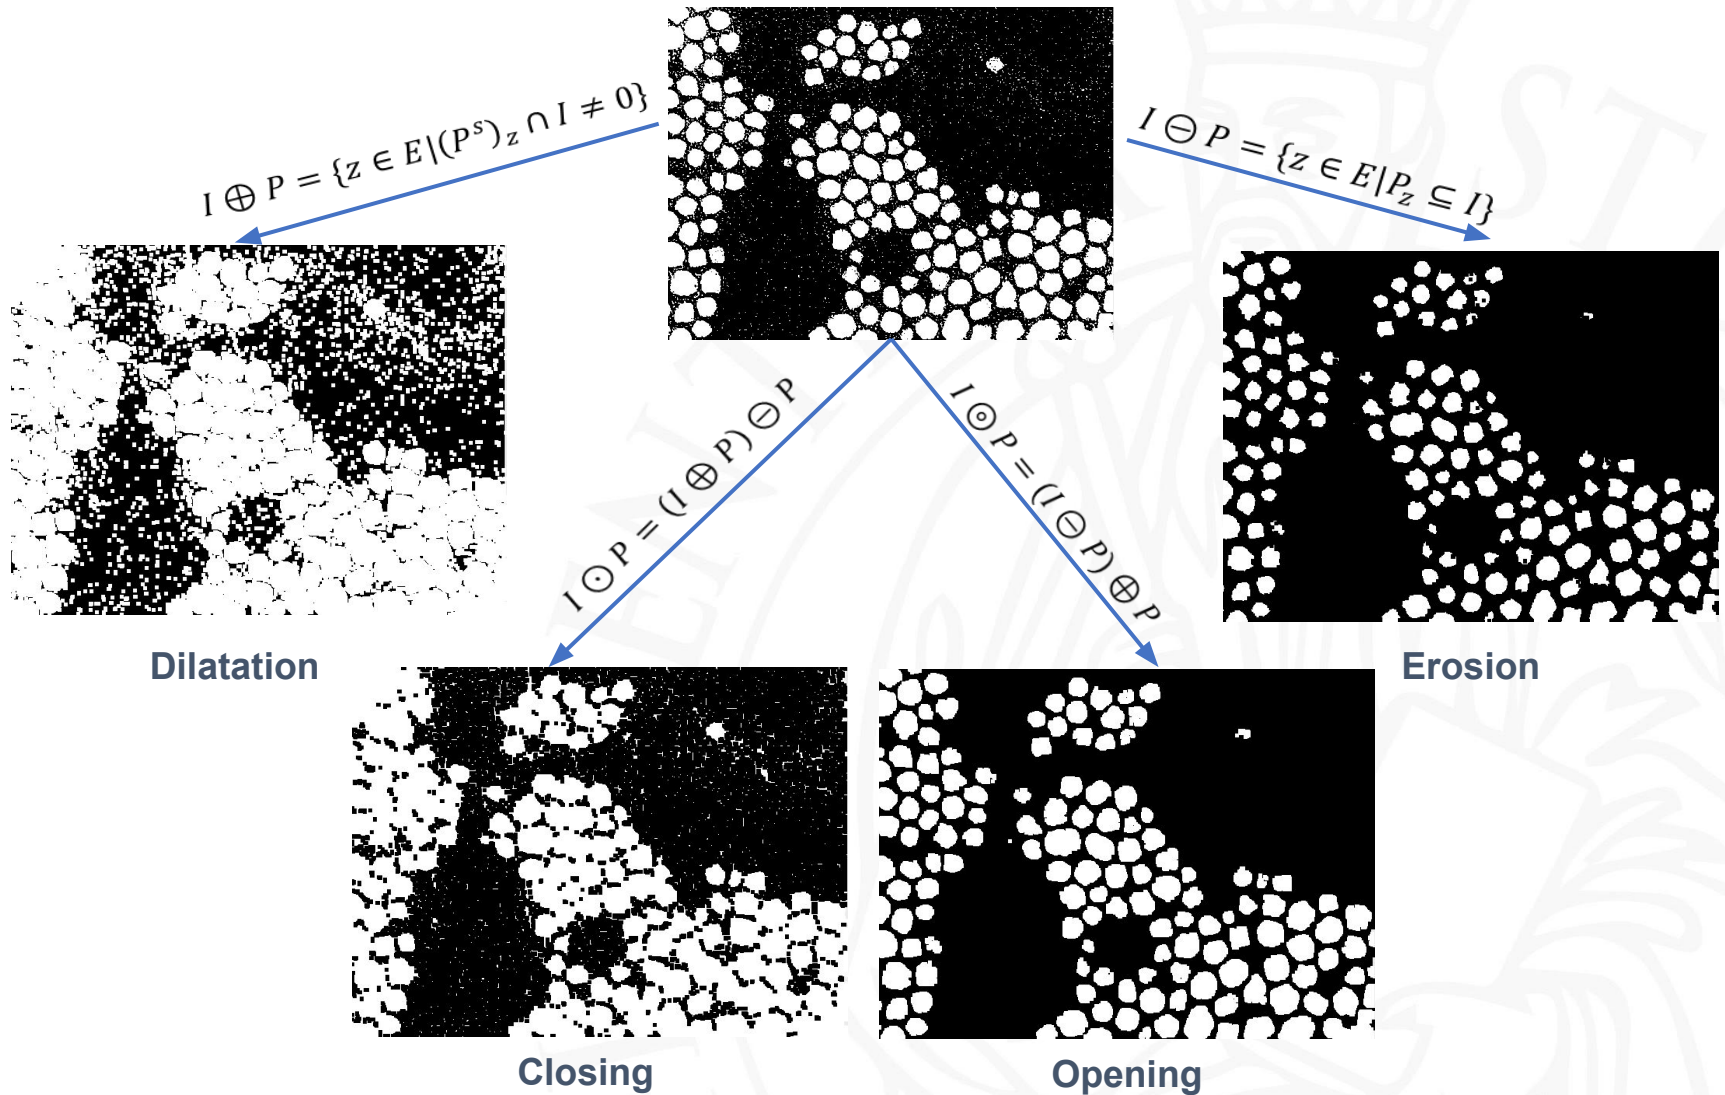
*Image (I):*



*Probe Elements (P):*



# Introduction: Morphological Operators




### Pseudo-Code Algorithm:

```
for rowImg = 0,1,...,imgH do  
  for colImg = 0,1,...,imgW do  
    for rowPrb = 0,1,...,probeH do  
      for colPrb = 0,1,...,probeW do  
        update Neighborhood[ ]  
      end for  
    end for  
    update Img[rowImg*imgW+colImg] with  
      max or min of Neighborhood  
  end for  
end for
```



# JAVA Implementation



- **ExecutorService** can be used as a **Fork-Join** pattern
- Easy to instantiate and execute a job
- Easy to create a **barrier** with **waitExecutors** method




## JAVA: parallel idea

Since computation of each pixel is **independent** from the others (apart from getting neighbours value) the idea is:

- Given a **Thread** number **N**
- Given a grayscale image with **size W, H**

Split the image in **N chunks of rows** and let each thread process his chunk of data.  
Finally keep a **barrier** and **wait** for all to complete.





## JAVA: implementation

(Dilation method)

```
int rowsOfTh = Math.floor(H/N); // max rows per thread
ExecutorService tasks_executor = Executors.newFixedThreadPool(N);
for (int r = 0; r < H; r += rowsOfTh){
    MaskApplier t = new MaskApplier(...);
    tasks_executor.execute(t)
}
waitExecutors(tasks_executor); //barrier method
```

(MaskApplier run method)

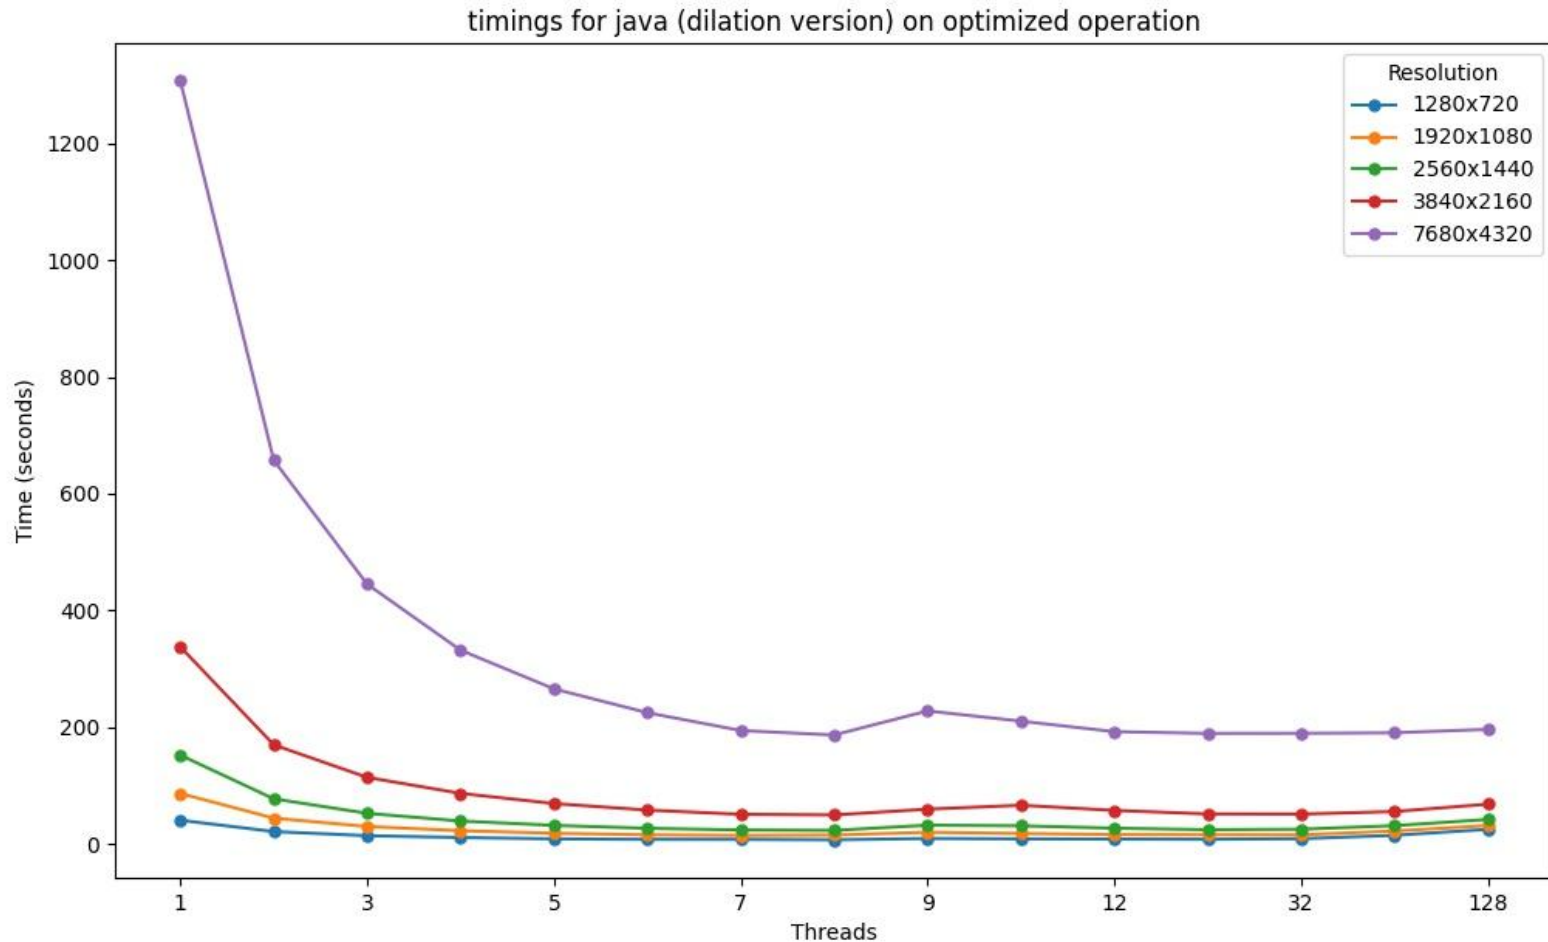
```
int m = MASK_SIZE;
for (y = y_min; y < y_max; y++){ // loop on Thread's rows
    for (x = 0; x < W; x++){ // loop on all x values
        p = img[x,y];
        for (ty = y-m/2; ty<=y+m/2; ty++)
            for (tx = x-m/2; tx<=x+m/2; tx++)
                if (p > img[tx, ty])
                    p = img[tx, ty] // find max on neighbours (mask)
        output_img[x, y] = p // set output pixel
    }
}
```

## JAVA: experiments



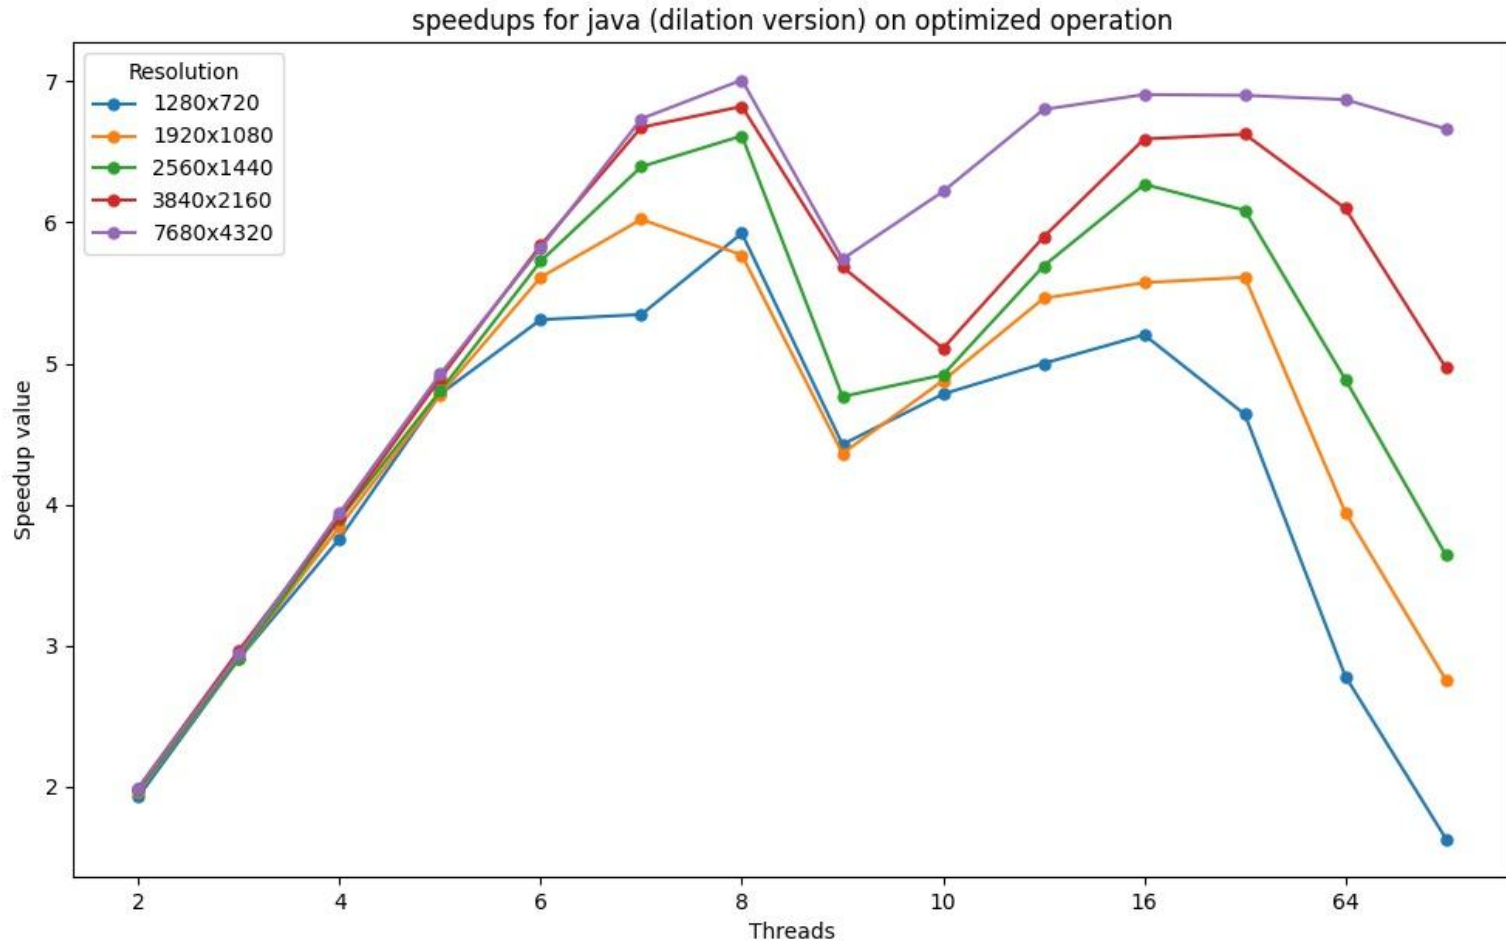
- **Sequential Timings** are taken by:  
CPU Intel(R) Core(TM) i7-9700K CPU @ 3.6GHz
- **Parallel Timings** are taken by:  
CPU Intel(R) Core(TM) i7-9700K CPU @ 3.6GHz  
Each core boosts up to 4.6 Ghz OOTB  
No HyperThreading
- **Experiments** with:
  - Dataset = {1280x720, 1920x1080, 2560x1440, 3840x2160, 7680x4320}
  - Threads = {2, 3, 4, 6, 8, ... , 128}

# JAVA: results (timing)



Results show a huge timing decrease until 8 Threads. Bigger datasets  $\Rightarrow$  bigger speedup.

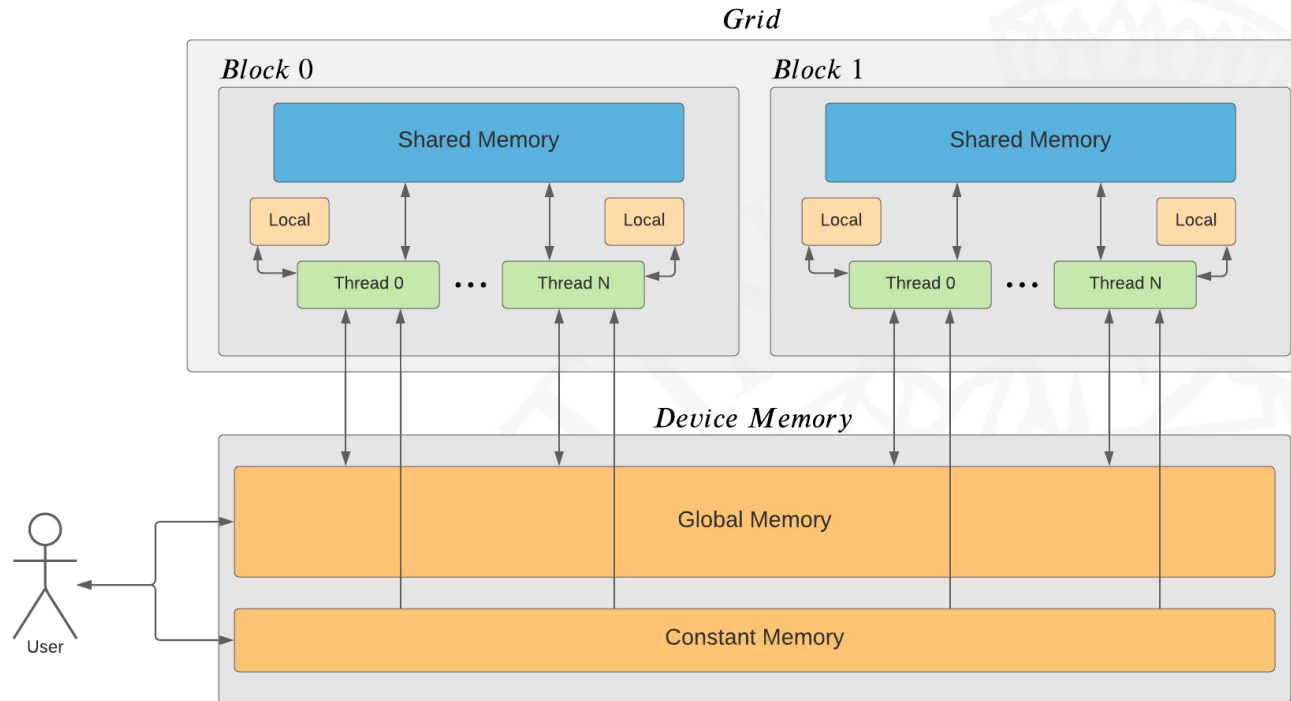
# JAVA: results (speedups)



Results show max speedup of 7 at 8 threads. Bigger datasets  $\Rightarrow$  bigger speedup.



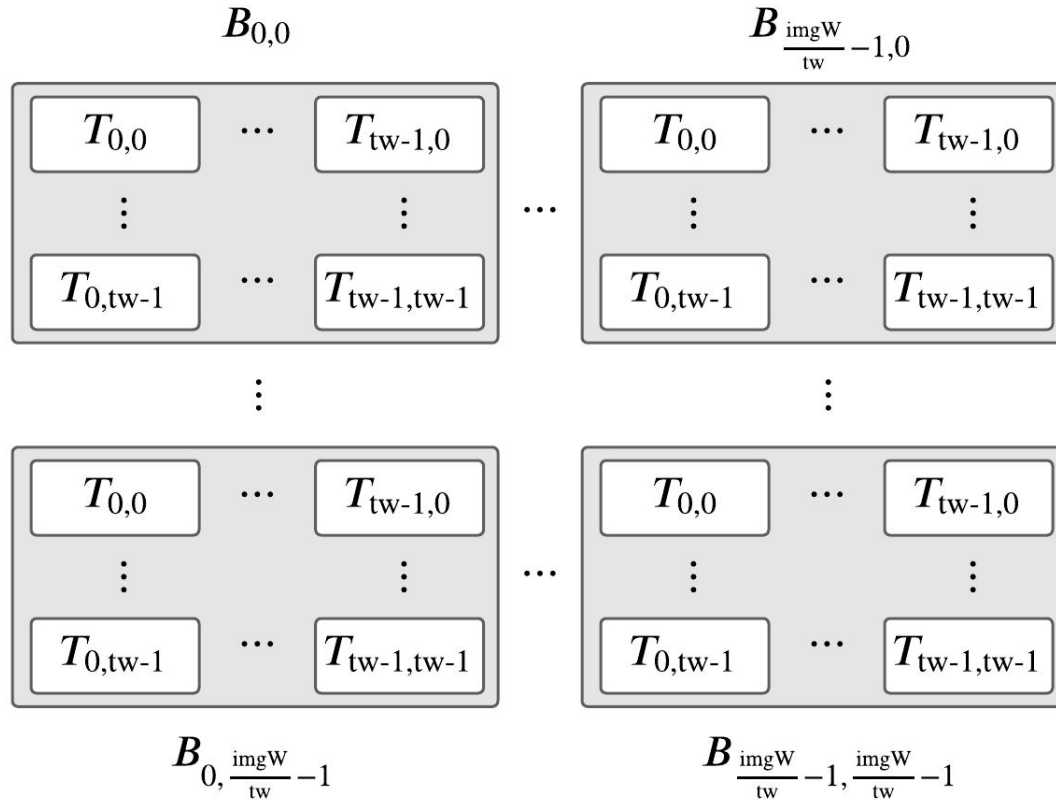
# CUDA Implementation



- **SIMT** (Single Instruction Multiple Thread) model.
- User decides number of threads  $T \rightarrow \left\lceil \frac{N}{T} \right\rceil$  blocks.
- Exploit both **global** and **shared** memory access.



# CUDA: Kernel Calls

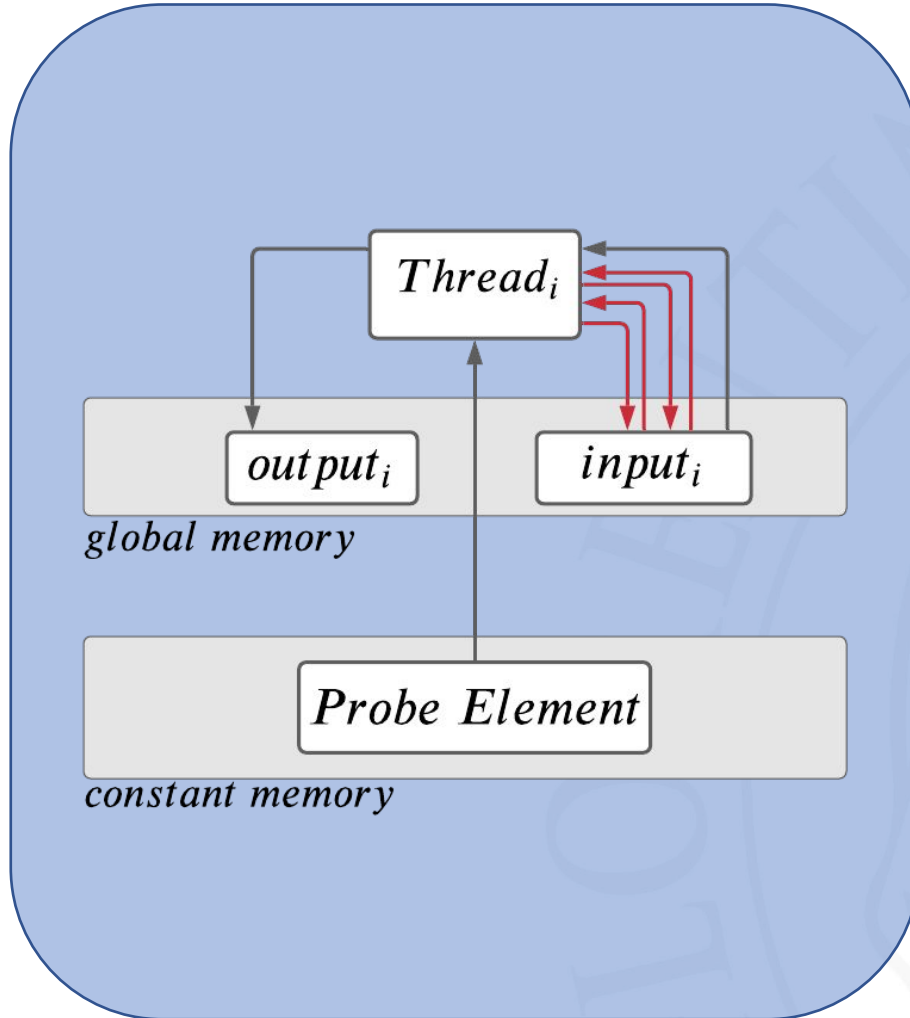


*Naive* → process  $\lll dimGrid, dimBlock \ggg (inImg, probe, outImg)$


*Optimized* → process  $\lll dimGrid, dimBlock, shared\_amount \ggg (inImg, probe, outImg)$

# CUDA: Solutions

Naive:



Optimized:



# CUDA: Naive Solution


## Pseudo-Code:

```

tid_x = (blockIdx.x * blockDim.x) + threadIdx.x
tid_y = (blockIdx.y * blockDim.y) + threadIdx.y
max = min = inImg[rowImg * imgW + colImg]
for rowPrb = 0, ... , prbH do
    for colPrb = 0, ... , prbW do
        update(max)
        update(min)
    end for
end for

If (EROSION) {outImg[rowImg * imgW + colImg] = min}
If (DILATATION) {outImg[rowImg * imgW + colImg] = max}


__syncthreads()
    
```



# CUDA: Optimized Solution

## 2-Batch Loading:

- $TW = \text{TILE\_WIDTH}$
- Each blocks is made of  $TW * TW$  threads.
- $W = TW + \text{Probe\_Width} - 1$



### First Batch Loading:

```
dest = threadIdx.y * TW + threadIdx.x
tile_DS[dest / W][dest % W] = inputImg[src]
```

### Second Batch Loading:

```
dest = threadIdx.y * TW + threadIdx.x + TW * TW
tile_DS[dest / W][dest % W] = inputImg[src]
```

```
__syncthreads()
```

```
for y = 0, ..., probeW do
  for x = 0, ..., probeH do
    compute max and min
  end for
end for
```

```
__syncthreads()
```


*shared memory*



*Thread<sub>i</sub>*



*global memory*



*constant memory*

# CUDA: Experiments



- **Sequential Timings** are taken by:

CPU Intel(R) Core(TM) i7-8750H CPU @ 2.20GHz



- **Parallel Timings** are taken by:

GPU NVIDIA GeForce GTX 1050 Ti

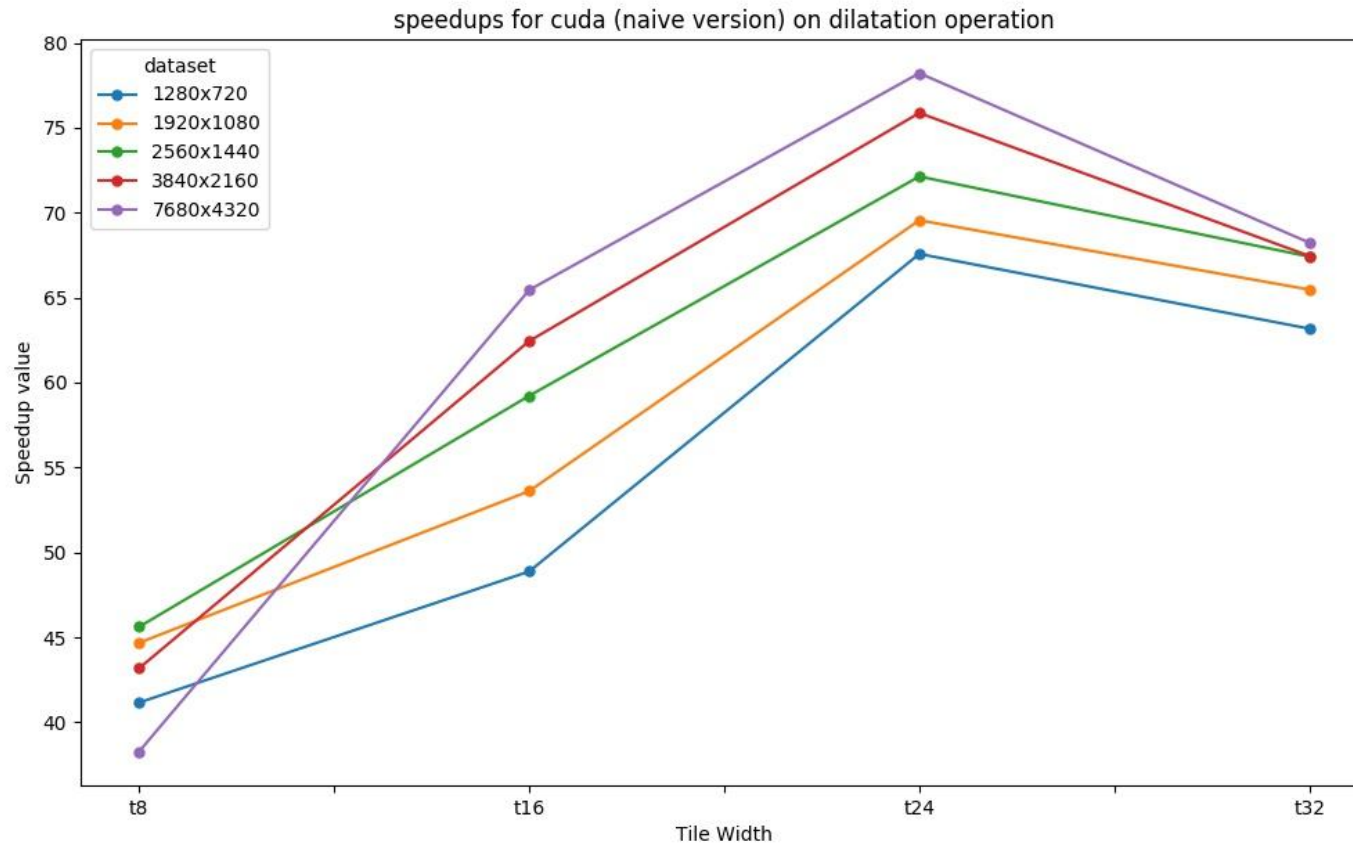
With 4096 MB dedicated and 768 CUDA Cores



- **Experiments** with:

- Dataset = {1280x720, 1920x1080, 2560x1440, 3840x2160, 7680x4320}
- Tiles = {8, 16, 24, 32}

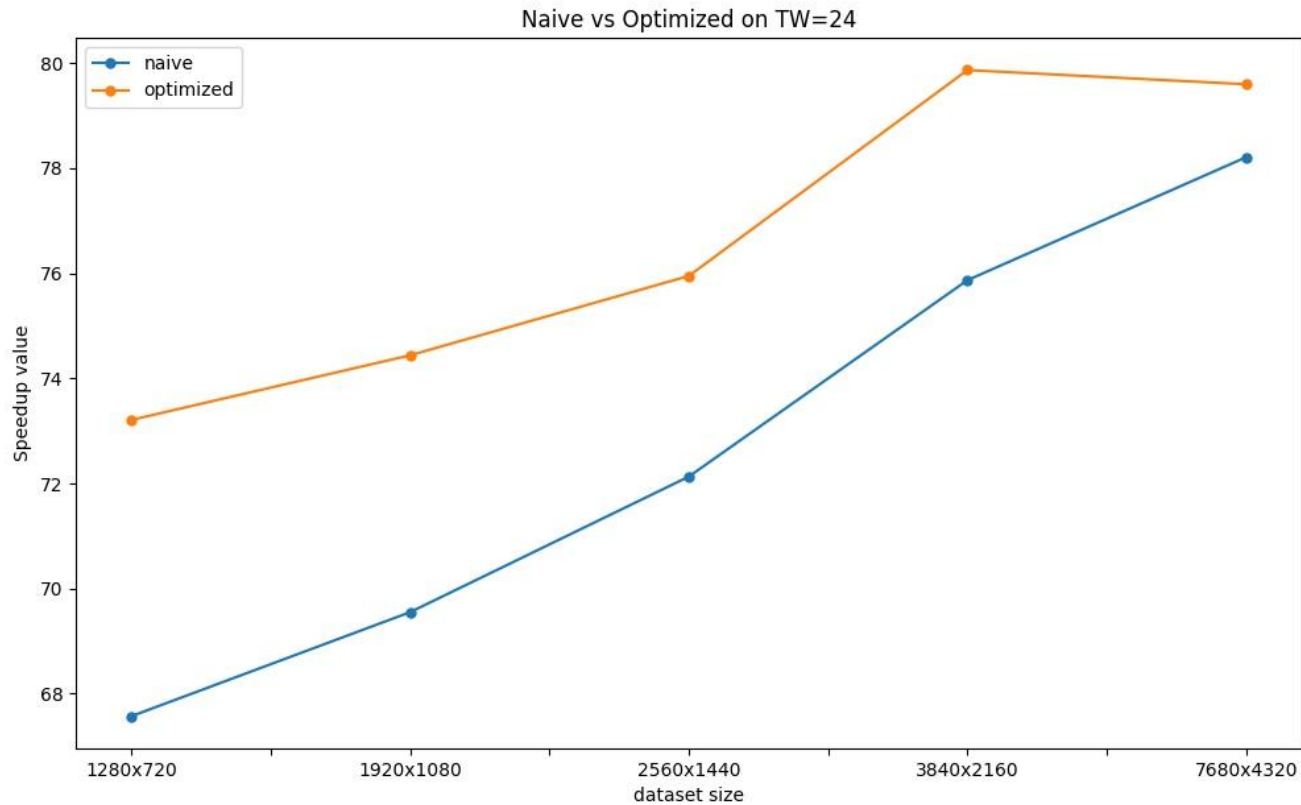
# CUDA: Results



- The best speedup is obtained with Tile Width = 24, thus 576 threads per block.
- Bigger Datasets  $\Rightarrow$  Bigger Speedups

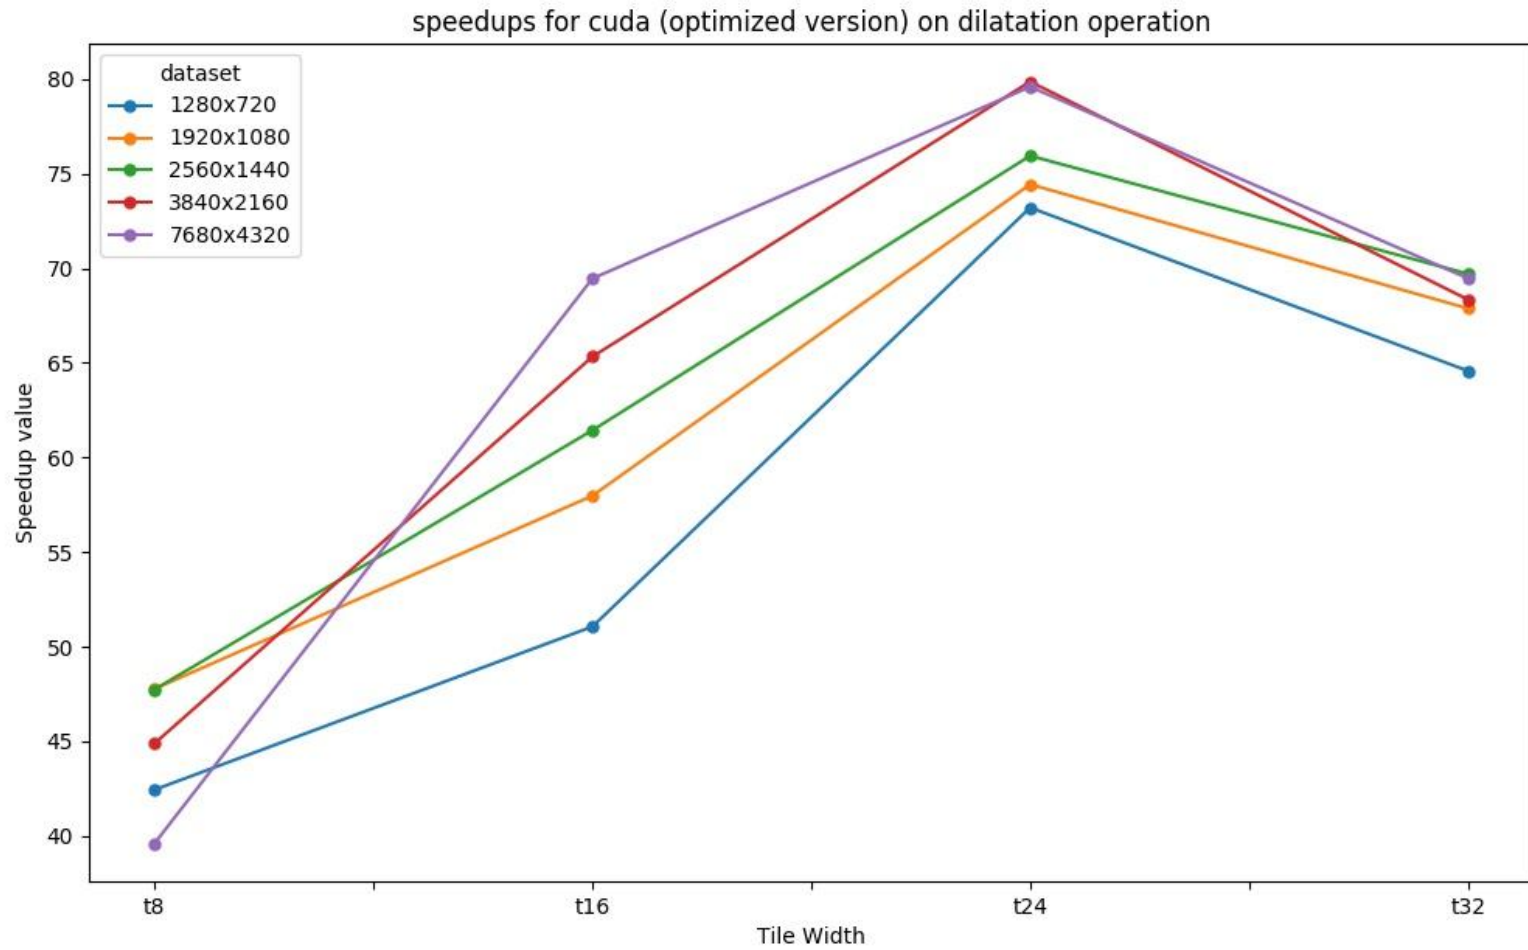


# CUDA: Results




- Max speedup = 91.09 is obtained with Tile Width = 24 and dataset = 3840x2160

# CUDA: Results



## Conclusions:

**Three implementations of MM: one sequential and two parallel.**



<https://github.com/AngeloDamante/morphological-image-processing-in-parallel>



UNIVERSITÀ  
DEGLI STUDI  
FIRENZE

# Thanks for your attention

Angelo D'Amante  
Fabian Greavu